

Name:

Jeffrey Hui

Netid:

jhui8

CS 441 - HW1: Instance-based Methods

Complete the sections below. You do not need to fill out the checklist.

Total Points Available

[] / 145

1. Retrieval, K-means, 1-NN on MNIST
 - a. Retrieval [] / 5
 - b. K-means [] / 15
 - c. 1-NN [] / 10
2. Make it fast
 - a. K-means plot [] / 15
 - b. 1-NN error plots [] / 8
 - c. 1-NN time plots [] / 7
 - d. Most confused label [] / 5
3. Temperature Regression
 - a. RMSE Tables [] / 20
4. Conceptual questions [] / 15
5. Stretch Goals
 - a. Evaluate effect of K for MNIST [] / 15
 - b. Evaluate effect of K for Temp Reg. [] / 15
 - c. Compare Kmeans more iterations vs. restarts [] / 15

1. Retrieval, K-means, 1-NN on MNIST

a. What index is returned for `x_test[1]`?

28882

b. Paste the display of clusters after the 1st and 10th iteration for $K=30$.


5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 2 8 6 9 4 0 9 1 1 2 9 3 2 7
 5 0 9 1 9 2 3 3 1 4 3 5 3 6 1 7 6 8 6 9 4 0 4 1 2 0 7 5 2 7

c. Error rate for first 100 test samples, using first 10,000 training samples (x.x)


8.0

2. Make it fast


a. KMeans plot of RMSE vs iterations for K=10, 30, 100



b. Nearest neighbor error vs training size plot



c. Nearest neighbor time vs training size plot



d. What label is most commonly confused with '2'?

7

3. Temperature Regression

a. Table of RMSE for KNN with K=5 (x.xx)

| | KNN (K=5) |
|---------------------|-----------|
| Original Features | 3.25 |
| Normalized Features | 2.93 |

4. Test your understanding

Fill in the letter corresponding to the answer. If you're not sure, you can sometimes run small experiments to check.

1. Is K-means guaranteed to decrease RMSE between nearest cluster and samples at each iteration until convergence?

- a. Yes
- b. No

b

2. If you increase K, is K-means expected or guaranteed to achieve lower RMSE?
- a. Guaranteed
 - b. Expected but not guaranteed
 - c. Not expected

b

3. In K-NN regression, for training labels y , what is the lowest target value that can possibly be predicted for any query?
- a. $\min(y)$
 - b. $\text{Mean}(y)$
 - c. Can't be determined

a

4. Would you expect the “training error” for 1-NN to be higher or lower than 3-NN for classification? Training error is the error if you test on the training data.
- a. Higher
 - b. Lower
 - c. It's problem-dependent

b

5. Would you expect the test error for 1-NN to be higher or lower than for 3-NN for regression?
- a. Higher
 - b. Lower
 - c. It's problem-dependent

a

5. Stretch Goals (optional)

- a. Select best K parameter for K-NN MNIST classification in K=1, 3, 5, 11, 25. (x.xx)

| Validation Set Performance | K=1 | K=3 | K=5 | K=11 | K=25 |
|----------------------------|-----|-----|-----|------|------|
|----------------------------|-----|-----|-----|------|------|

| | | | | | |
|---------|------|------|------|------|------|
| % error | 2.88 | 2.80 | 2.82 | 3.08 | 3.82 |
|---------|------|------|------|------|------|

Best K:

3

Test % error (x.xx)

2.95

b. Select best K parameter for K-NN temperature regression in K=1, 3, 5, 11, 25. (x.xx)

| Validation Set RMSE | K=1 | K=3 | K=5 | K=11 | K=25 |
|---------------------|-----|-----|-----|------|------|
| Original Features | | | | | |
| Normalized Features | | | | | |

Best Setting (K, feature type):

Test RMSE (x.xx)

c. Kmeans, MNIST: compare average and standard deviation RMSE based on number of iterations and number of restarts

(4 digit precision)

| K=30 | RMSE avg | RMSE std |
|---------------------------|----------|----------|
| 20 iterations, 1 restart | 5.7912 | 0.0114 |
| 4 iterations, 5 restarts | 5.8302 | 0.0115 |
| 50 iterations, 1 restart | 5.7766 | 0.0029 |
| 10 iterations, 5 restarts | 5.7894 | 0.0099 |

Acknowledgments / Attribution

List any outside sources for code or ideas or “None”.

None.

cs441-sp24-hw1-starter

February 6, 2024

0.1 CS441: Applied ML - HW 1

0.1.1 Parts 1-2: MNIST

Include all the code for generating MNIST results below

```
[1]: # initialization code
import numpy as np
from keras.datasets import mnist
%matplotlib inline
from matplotlib import pyplot as plt
from scipy import stats

def load_mnist():
    '''
    Loads, reshapes, and normalizes the data
    '''
    (x_train, y_train), (x_test, y_test) = mnist.load_data() # loads MNIST data
    x_train = np.reshape(x_train, (len(x_train), 28*28)) # reformat to 768-d
    ↪vectors
    x_test = np.reshape(x_test, (len(x_test), 28*28))
    maxval = x_train.max()
    x_train = x_train/maxval # normalize values to range from 0 to 1
    x_test = x_test/maxval
    return (x_train, y_train), (x_test, y_test)

def display_mnist(x, subplot_rows=1, subplot_cols=1):
    '''
    Displays one or more examples in a row or a grid
    '''
    if subplot_rows>1 or subplot_cols>1:
        fig, ax = plt.subplots(subplot_rows, subplot_cols, figsize=(15,15))
        for i in np.arange(len(x)):
            ax[i].imshow(np.reshape(x[i], (28,28)), cmap='gray')
            ax[i].axis('off')
    else:
        plt.imshow(np.reshape(x, (28,28)), cmap='gray')
        plt.axis('off')
```



```
plt.show()
```

WARNING:tensorflow:From D:\CODE\School\cs441\venv\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
[2]: # example of using MNIST load, display, indices, and count functions
(x_train, y_train), (x_test, y_test) = load_mnist()
display_mnist(x_train[:10], 1, 10)
print('Total size: train={}, test {}'.format(len(x_train), len(x_test)))
```



Total size: train=60000, test =10000

1. Retrieval, Clustering, and NN Classification

```
[3]: # Retrieval

def get_nearest(X_query, X):
    ''' Return the index of the sample in X that is closest to X_query according
        to L2 distance '''
    return np.argmin([np.sum((X_query - x) ** 2) for x in X])

j = get_nearest(x_test[0], x_train)
print(j)
j = get_nearest(x_test[1], x_train)
print(j)
```

53843

28882

```
[4]: # K-means
import copy
def kmeans(X, K, niter=10):
    '''
    Starting with the first K samples in X as cluster centers, iteratively assign
    each
    point to the nearest cluster and compute the mean of each cluster.
    Input: X[i] is the ith sample, K is the number of clusters, niter is the
    number of iterations
    Output: K cluster centers
```

```
'''
# TO DO -- add code to display cluster centers at each iteration also
clusters = [[] for _ in range(K)]
cluster_centers = copy.deepcopy(X[:K])

for i in range(niter):

    # For each x_i, compute the closest cluster it belongs
    for x_i in X:
        nearest_center = get_nearest(x_i, cluster_centers)
        clusters[nearest_center].append(x_i)
    # Update centers
    for j in range(len(cluster_centers)):
        cluster_centers[j] = np.mean(clusters[j], axis=0)

    display_mnist(cluster_centers, 1, K)
    # reset clusters
    clusters = [[] for _ in range(K)]

return cluster_centers
```

K=30

centers = kmeans(x_train[:1000], K)

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 2 9 3 2 9

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 9

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 9

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 9

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 9

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 9

509192331435361708694041207529

509192331435361708694041207529

509192331435361708694041207529

509192331435361708694041207529

```
[5]: # 1-NN
def one_nn(X_test, Y_test, X_train, Y_train):
    correct_labels = Y_test

    predictions = [get_nearest(x_t, X_train) for x_t in X_test]
    predictions = np.take(Y_train, predictions, axis=0)

    accuracy = (predictions == correct_labels).sum() / len(X_test)

    return 1 - accuracy

# TO DO
n_test = 100
n_train = 10000
error_rate = one_nn(x_test[:n_test], y_test[:n_test], x_train[:n_train],
    ↪ y_train[:n_train])
print('Error rate: {:.5f}'.format(error_rate))
```

Error rate: 0.08000

2. Make it fast

```
[6]: # install libraries you need for part 2
import faiss
import time
```

```
[7]: # retrieval

# TO DO (check that you're using FAISS correctly)
def get_nearest_faiss(X_query, X):
    index = faiss.IndexFlatL2(X.shape[1])
    index.add(X)
    dist, idx = index.search(X_query,1)

    return dist.item(), idx.item()

j = get_nearest_faiss(x_test[0].reshape(1,-1), x_train)
print(j)
j = get_nearest_faiss(x_test[1].reshape(1,-1), x_train)
print(j)
```

```
(7.039846420288086, 53843)
(20.798309326171875, 28882)
```

```
[8]: # K-means
from collections import defaultdict
def create_index(index_dataset):
    index = faiss.IndexFlatL2(index_dataset.shape[1])
    index.add(index_dataset)
    return index

def find_with_index(index: faiss.IndexFlatL2, query):
    dist, idx = index.search(query,1)
    return dist, idx

def kmeans_fast(X, K, niter=10):
    '''
    Starting with the first K samples in X as cluster centers, iteratively assign
    each
    point to the nearest cluster using faiss and compute the mean of each cluster.
    Input: X[i] is the ith sample, K is the number of clusters, niter is the
    number of iterations
    Output: K cluster centers
    '''

    # TO DO (you can base this on part 1, but use FAISS for search)
    # if you include display code, you need to re-organize the plotting code below

    clusters = defaultdict(list)
    distances = []
    cluster_centers = copy.deepcopy(X[:K])

    for i in range(niter):
```

```

faiss_index = create_index(cluster_centers)
dist, nearest_center = find_with_index(faiss_index, X)
dist, nearest_center = dist.flatten(), nearest_center.flatten()
distances.append(np.sqrt(np.mean(dist)))

# Populate clusters
for img_idx, center_idx in enumerate(nearest_center):
    clusters[center_idx].append(X[img_idx])

# Update centers
for j in range(len(cluster_centers)):
    cluster_centers[j] = np.mean(clusters[j], axis=0)

# reset
clusters = defaultdict(list)


return cluster_centers, distances

K=10
centers, rmse = kmeans_fast(x_train, K, niter=20)
plt.plot(np.arange(len(rmse)), rmse, label='K=10')

K=30
centers, rmse = kmeans_fast(x_train, K, niter=20)
plt.plot(np.arange(len(rmse)), rmse, label='K=30')

K=100
centers, rmse = kmeans_fast(x_train, K, niter=20)
plt.plot(np.arange(len(rmse)), rmse, label='K=100')
plt.legend(), plt.ylabel('RMSE'), plt.xlabel('# iterations')
plt.show()

```



```
[9]: # 1-NN
from time import time
nsample = [100, 1000, 10000, 60000]

calc_acc = lambda actual, prediction: (prediction == actual).sum() / len(prediction)
get_predictions = lambda indices: np.take(y_train, indices, axis=0)
acc = lambda actual, prediction: calc_acc(actual, get_predictions(prediction))

acc_exact = []
acc_lsh = []

timing_exact = []
timing_lsh = []

exact_predictions = None

# TO DO
for samp in nsample:
    # setup index for exact
    dim = x_train.shape[1]
```

```

exact_idx = faiss.IndexFlatL2(dim)
exact_idx.add(x_train[:samp])

# search for exact
t1 = time()
_, idx = exact_idx.search(x_test, 1) # full test set search
t2 = time()
timing_exact.append(t2 - t1)
acc_exact.append(acc(y_test, idx.reshape(-1))) # take out the extra dimension
↳ using reshape
exact_predictions = get_predictions(idx)

# setup index for lsh
lsh_idx = faiss.IndexLSH(dim, dim)
lsh_idx.add(x_train[:samp])


# search for lsh
t1 = time()
_, idx = lsh_idx.search(x_test, 1)
t2 = time()
timing_lsh.append(t2 - t1)
acc_lsh.append(acc(y_test, idx.reshape(-1)))


acc_exact = np.array(acc_exact)
acc_lsh = np.array(acc_lsh)

plt.semilogx(nsamples, 1-acc_exact, label='exact')
plt.semilogx(nsamples, 1-acc_lsh, label='lsh')
plt.legend(), plt.ylabel('% error'), plt.xlabel('# training samples')
plt.show()

plt.semilogx(nsamples, timing_exact, label='exact')
plt.semilogx(nsamples, timing_lsh, label='lsh')
plt.legend(), plt.ylabel('time'), plt.xlabel('# training samples')
plt.show()

```






```
[10]: # Confusion matrix
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

cm = confusion_matrix(y_test, exact_predictions)
disp = ConfusionMatrixDisplay(cm)
disp.plot()

# TO DO
```

```
[10]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x278f2054cd0>
```



0.2 Part 3: Temperature Regression

Include all your code used for part 2 in this section.

```
[11]: import numpy as np
      %matplotlib inline
      from matplotlib import pyplot as plt
      from sklearn.linear_model import Ridge
      from sklearn.linear_model import Lasso

      # load data (modify to match your data directory or comment)
      def load_temp_data():
          datadir = "./"
          T = np.load(datadir + 'temperature_data.npz')
          x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val, \
          ↪ dates_test, feature_to_city, feature_to_day = \
          T['x_train'], T['y_train'], T['x_val'], T['y_val'], T['x_test'], T['y_test'], \
          ↪ T['dates_train'], T['dates_val'], T['dates_test'], T['feature_to_city'], \
          ↪ T['feature_to_day']
          return (x_train, y_train, x_val, y_val, x_test, y_test, dates_train, \
          ↪ dates_val, dates_test, feature_to_city, feature_to_day)
```

```

# plot one data point for listed cities and target date
def plot_temps(x, y, cities, feature_to_city, feature_to_day, target_date):
    nc = len(cities)
    ndays = 5
    xplot = np.array([-5,-4,-3,-2,-1])
    yplot = np.zeros((nc,ndays))
    for f in np.arange(len(x)):
        for c in np.arange(nc):
            if cities[c]==feature_to_city[f]:
                yplot[feature_to_day[f]+ndays,c] = x[f]
    plt.plot(xplot,yplot)
    plt.legend(cities)
    plt.plot(0, y, 'b*', markersize=10)
    plt.title('Predict Temp for Cleveland on ' + target_date)
    plt.xlabel('Day')
    plt.ylabel('Avg Temp (C)')
    plt.show()

```

```

[12]: # load data
(x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val,
  ↪ dates_test, feature_to_city, feature_to_day) = load_temp_data()
''' Data format:
    x_train, y_train: features and target value for each training sample
    ↪ (used to fit model)
    x_val, y_val: features and target value for each validation sample (used
    ↪ to select hyperparameters, such as regularization and K)
    x_test, y_test: features and target value for each test sample (used to
    ↪ evaluate final performance)
    dates_xxx: date of the target value for the corresponding sample
    feature_to_city: maps from a feature number to the city
    feature_to_day: maps from a feature number to a day relative to the
    ↪ target value, e.g. -2 means two days before
    Note: 361 is the temperature of Cleveland on the previous day
'''
f = 361
print('Feature {}: city = {}, day= {}'.format(f,feature_to_city[f],
  ↪ feature_to_day[f]))
baseline_rmse = np.sqrt(np.mean((y_val[1:]-y_val[:-1])**2)) # root mean squared
  ↪ error example
print('Baseline - prediction using previous day: RMSE={}'.format(baseline_rmse))


# plot first two x/y for val
plot_temps(x_val[0], y_val[0], ['Cleveland', 'New York', 'Chicago', 'Denver',
  ↪ 'St. Louis'], feature_to_city, feature_to_day, dates_val[0])


```

```
plot_temps(x_val[1], y_val[1], ['Cleveland', 'New York', 'Chicago', 'Denver', 'St. Louis'], feature_to_city, feature_to_day, dates_val[1])
```

Feature 361: city = Cleveland, day= -1

Baseline - prediction using previous day: RMSE=3.460601246750482





```
[13]: # K-NN Regression

def regress_KNN(X_trn, y_trn, X_tst, K=1):
    '''
    Predict the target value for each data point in X_tst using a
    K-nearest neighbor regressor based on (X_trn, y_trn), with L2 distance.
    Input: X_trn[i] is the ith training data. y_trn[i] is the ith training label.
    ↪ K is the number of closest neighbors to use.
    Output: return y_pred, where y_pred[i] is the predicted ith test value
    '''
    # TO DO
    faiss_index = create_index(X_trn)
    dist, idx = faiss_index.search(X_tst, K)

    # get prediction
    predictions = np.take(y_trn, idx, axis=0)
    predictions = np.mean(predictions, axis=1)

    return predictions
```

```

def normalize_features(x, y, fnum):
    ''' Normalize the features in x and y.
        For each data sample i:
            x2[i] = x[i]-x[i,fnum]
            y2[i] = y[i]-x[i,fnum]
    '''
    # TO DO
    return (x - np.take(x, fnum, axis=1).reshape(-1, 1)), (y - np.take(x, fnum,
↵axis=1))

# KNN with original features
res = regress_KNN(x_train, y_train, x_test, 5)
print(np.sqrt(np.mean((res - y_test) ** 2)))

# TO DO

# KNN with normalized features
fnum = 361 # previous day temp in Cleveland

# TO DO
x_train_norm, y_train_norm = normalize_features(x_train, y_train, fnum)
x_test_norm, y_test_norm = normalize_features(x_test, y_test, fnum)
# KNN with normalized features
res = regress_KNN(x_train_norm, y_train_norm, x_test_norm, 5)
print(np.sqrt(np.mean((res - y_test_norm) ** 2)))

```

3.249556245363484
2.9324389176041588

0.3 Part 5: Stretch Goals

Include all your code used for part 5 in this section. You can copy-paste code from parts 1-3 if it is re-usable.

```

[14]: # Stretch: KNN classification (Select K)
from sklearn.neighbors import KNeighborsClassifier

(x_train, y_train), (x_test, y_test) = load_mnist()

p5a_train_x, p5a_train_y = x_train[:50000], y_train[:50000]
p5a_val_x, p5a_val_y = x_train[50000:], y_train[50000:]

# for i in [1,3,5,11,25]:
#     neigh = KNeighborsClassifier(n_neighbors=i)
#     neigh.fit(p5a_train_x, p5a_train_y)
#     print(f'n = {i}, error rate = {1 - neigh.score(p5a_val_x, p5a_val_y)}')

```

```

neigh = KNeighborsClassifier(n_neighbors=3)
neigh.fit(x_train, y_train)
1 - neigh.score(x_test, y_test)

```

[14]: 0.029499999999999997

[15]: *# Stretch: KNN regression (Select K)*

[16]: *# Stretch: K-means (more iters vs redos*

```

params = [(20, 1), (4, 5), (50, 1), (10, 5)]

for (iteration, restart) in params:
    rmses = []
    for _ in range(5):
        fkmeans = faiss.Kmeans(x_train.shape[1], 30, niter=iteration,
↪nredo=restart, seed=int(time()))
        fkmeans.train(x_train)
        dist, idx = fkmeans.index.search(x_train, 1)
        rmses.append(np.sqrt(np.sum(dist) / x_train.shape[0]))

    std = np.std(rmses)
    avg = np.mean(rmses)

    print(f'iteration: {iteration}, redos: {restart}, mean: {avg}, std: {std}')

```

```

iteration: 20, redos: 1, mean: 5.79427435672978, std: 0.013635576397397473
iteration: 4, redos: 5, mean: 5.824453619315952, std: 0.0014559432134116084
iteration: 50, redos: 1, mean: 5.780731385540949, std: 0.008345551020031898
iteration: 10, redos: 5, mean: 5.780008829446705, std: 0.003928297103671805

```